

Expert Oracle Database 11*g* Administration



Sam R. Alapati

Apress®

Expert Oracle Database 11g Administration

Copyright © 2009 by Sam R. Alapati

All rights reserved. No part of this work may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval system, without the prior written permission of the copyright owner and the publisher.

ISBN-13 (pbk): 978-1-4302-1015-3

ISBN-13 (electronic): 978-1-4302-1016-0

Printed and bound in the United States of America 9 8 7 6 5 4 3 2 1

Trademarked names may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, we use the names only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

Lead Editor: Jonathan Gennick

Developmental Editor: Douglas Pundick

Technical Reviewer: John Watson

Editorial Board: Clay Andres, Steve Anglin, Mark Beckner, Ewan Buckingham, Tony Campbell,

Gary Cornell, Jonathan Gennick, Michelle Lowman, Matthew Moodie, Jeffrey Pepper,

Frank Pohlmann, Ben Renow-Clarke, Dominic Shakeshaft, Matt Wade, Tom Welsh

Project Manager: Richard Dal Porto

Copy Editor: Ami Knox

Associate Production Director: Kari Brooks-Copony

Production Editor: Laura Cheu

Compositors: Susan Glinert Stevens, Ellie Fountain

Proofreader: April Eddy

Indexer: John Collin

Cover Designer: Kurt Krames

Manufacturing Director: Tom Debolski

Distributed to the book trade worldwide by Springer-Verlag New York, Inc., 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax 201-348-4505, e-mail orders-ny@springer-sbm.com, or visit <http://www.springeronline.com>.

For information on translations, please contact Apress directly at 2855 Telegraph Avenue, Suite 600, Berkeley, CA 94705. Phone 510-549-5930, fax 510-549-5939, e-mail info@apress.com, or visit <http://www.apress.com>.

Apress and friends of ED books may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Special Bulk Sales—eBook Licensing web page at <http://www.apress.com/info/bulksales>.

The information in this book is distributed on an “as is” basis, without warranty. Although every precaution has been taken in the preparation of this work, neither the author(s) nor Apress shall have any liability to any person or entity with respect to any loss or damage caused or alleged to be caused directly or indirectly by the information contained in this work.

To my dear wife, Valerie, as a small token for her immense help and support

Contents at a Glance

| | |
|------------------------------------|--------|
| About the Author | xxxvii |
| About the Technical Reviewer | xxxix |
| Acknowledgments | xli |
| Introduction | xlili |

PART 1 ■■■ Background, Data Modeling, UNIX/Linux, and SQL*Plus

| | | |
|-------------|--|----|
| ■ CHAPTER 1 | The Oracle DBA's World | 3 |
| ■ CHAPTER 2 | Relational Database Modeling and Database Design | 19 |
| ■ CHAPTER 3 | Essential UNIX (and Linux) for the Oracle DBA | 43 |
| ■ CHAPTER 4 | Using SQL*Plus and Oracle Enterprise Manager | 97 |

PART 2 ■■■ Oracle Database 11*g* Architecture, Schema, and Transaction Management

| | | |
|-------------|--|-----|
| ■ CHAPTER 5 | Oracle Database 11 <i>g</i> Architecture | 165 |
| ■ CHAPTER 6 | Managing Tablespaces | 215 |
| ■ CHAPTER 7 | Schema Management | 261 |
| ■ CHAPTER 8 | Oracle Transaction Management | 337 |

PART 3 ■■■ Installing Oracle Database 11*g*, Upgrading, and Creating Databases

| | | |
|--------------|---|-----|
| ■ CHAPTER 9 | Installing and Upgrading to Oracle Database 11 <i>g</i> | 391 |
| ■ CHAPTER 10 | Creating a Database | 443 |

PART 4 ■■■ Connectivity and User Management

| | | |
|--------------|---|-----|
| ■ CHAPTER 11 | Connectivity and Networking | 511 |
| ■ CHAPTER 12 | User Management and Database Security | 543 |

PART 5 ■■■ Data Loading, Backup, and Recovery

| | | |
|--------------|---|-----|
| ■ CHAPTER 13 | Loading and Transforming Data | 625 |
| ■ CHAPTER 14 | Using Data Pump Export and Import | 677 |
| ■ CHAPTER 15 | Backing Up Databases | 725 |
| ■ CHAPTER 16 | Database Recovery | 801 |

PART 6 ■■■ Managing the Database

| | | |
|--------------|--|-----|
| ■ CHAPTER 17 | Automatic Management and Online Capabilities | 877 |
| ■ CHAPTER 18 | Managing and Monitoring the Operational Database | 947 |

PART 7 ■■■ Performance Tuning

| | | |
|--------------|--|------|
| ■ CHAPTER 19 | Improving Database Performance: SQL Query Optimization ... | 1041 |
| ■ CHAPTER 20 | Performance Tuning: Tuning the Instance | 1129 |
| ■ APPENDIX | Oracle Database 11g SQL and PL/SQL: A Brief Primer | 1221 |
| ■ INDEX | | 1255 |

Contents

About the Author xxxvii

About the Technical Reviewer xxxix

Acknowledgments xli

Introduction xliii

PART 1 ■■■ **Background, Data Modeling, UNIX/
Linux, and SQL*Plus**

■ **CHAPTER 1 The Oracle DBA's World 3**

 The Oracle DBA's Role 3

 The DBA's Security Role 4

 The DBA's System Management Role 5

 The DBA's Database Design Role 7

 DBA Job Classifications 8

 Types of Databases 9

 Online Transaction Processing and Decision-Support

 System Databases 9

 Development, Test, and Production Databases 9

 Training and Certification 10

 Training 10

 Certification 11

 Resources and Organizations for Oracle DBAs 13

 Oracle by Example 14

 Oracle Database Two-Day DBA Course 14

 Oracle MetaLink 15

 Oracle Web Conference 15

 The Daily Routine of a Typical Oracle DBA 15

 Some General Advice 16

 Know When You Need Help 16

 Remember You Are Not Alone 16

 Think Outside the Box 17

 Primum Non Nocere 17

| | | |
|------------------|---|-----------|
| CHAPTER 2 | Relational Database Modeling and Database Design | 19 |
| | Relational Databases: A Brief Introduction | 19 |
| | The Relational Database Model | 20 |
| | Database Schemas | 20 |
| | Relational Algebra | 21 |
| | Relational Calculus | 22 |
| | SQL | 22 |
| | Relational Database Life Cycle | 23 |
| | Requirements Gathering and Analysis | 23 |
| | Logical Database Design | 24 |
| | Physical Database Design | 34 |
| | Implementing the Physical Design | 37 |
| | Reverse-Engineering a Database | 38 |
| | Object-Relational and Object Databases | 38 |
| | The Relational Model | 38 |
| | The Object Model | 39 |
| | The Object-Relational Model | 39 |
| | Semi-Structured Data Models | 40 |
| CHAPTER 3 | Essential UNIX (and Linux) for the Oracle DBA | 43 |
| | Overview of UNIX and Linux Operating Systems | 43 |
| | UNIX | 43 |
| | Linux | 44 |
| | Midrange Systems | 45 |
| | Understanding the UNIX Shell(s) | 45 |
| | Accessing the UNIX System | 46 |
| | Overview of Basic UNIX Commands | 48 |
| | Getting Help: The man Command | 50 |
| | Changing the Prompt | 51 |
| | Finding Files and Directories | 52 |
| | Controlling the Output of Commands | 52 |
| | Showing the Contents of Files | 52 |
| | Comparing Files | 53 |
| | Understanding Operating-System and Shell Variables | 53 |
| | Redirecting Input and Output | 56 |
| | Protecting Files from Being Overwritten | 57 |
| | Navigating Files and Directories in UNIX | 57 |
| | Files in the UNIX System | 57 |
| | Linking Files | 57 |

| | |
|---|----|
| Managing Files | 58 |
| Permissions: Reading from or Writing to Files in UNIX | 59 |
| Directory Management | 62 |
| Editing Files with the vi Editor | 63 |
| Creating and Modifying Files Using vi | 63 |
| Moving Around with the head and tail Commands | 65 |
| Extracting and Sorting Text | 65 |
| Using grep to Match Patterns | 65 |
| Cutting, Pasting, and Joining Text | 66 |
| Shell Scripting | 68 |
| What Is a Shell Program? | 68 |
| Using Shell Variables | 69 |
| Evaluating Expressions with the test Command | 69 |
| Executing Shell Programs with Command-Line Arguments | 70 |
| Analyzing a Shell Script | 70 |
| Flow-Control Structures in Korn Shell Programming | 71 |
| Dealing with UNIX Processes | 74 |
| Gathering Process Information with ps | 74 |
| Running Processes after Logging Out | 75 |
| Running Processes in the Background | 75 |
| Terminating Processes with the kill Command | 75 |
| UNIX System Administration and the Oracle DBA | 76 |
| UNIX Backup and Restore Utilities | 76 |
| The crontab and Automating Scripts | 77 |
| Using Telnet | 78 |
| Remote Login and Remote Copy | 78 |
| Using SSH, the Secure Shell | 79 |
| Using FTP to Send and Receive Files | 79 |
| UNIX System Performance Monitoring Tools | 80 |
| Disks and Storage in UNIX | 85 |
| Disk Storage Configuration Choices | 86 |
| Monitoring Disk Usage | 86 |
| Disk Storage, Performance, and Availability | 87 |
| RAID Systems | 88 |
| RAID Levels | 89 |
| Choosing the Ideal Disk Configuration | 90 |
| Redundant Disk Controllers | 92 |
| RAID and Backups | 92 |
| RAID and Oracle | 93 |

| | |
|---|----|
| Other Storage Technologies | 93 |
| Storage Area Networks | 93 |
| Networked Attached Storage..... | 94 |
| InfiniBand..... | 94 |
| Automatic Storage Management | 94 |
| Oracle and Storage System Compatibility | 95 |

CHAPTER 4 Using SQL*Plus and Oracle Enterprise Manager 97

| | |
|---|-----|
| Starting a SQL*Plus Session | 97 |
| Setting the Environment..... | 98 |
| Starting a SQL*Plus Session from the Command Line..... | 98 |
| Connecting by Using the CONNECT Command | 100 |
| Connectionless SQL*Plus Session with /NOLOG..... | 101 |
| Connecting to SQL*Plus Through a Windows GUI..... | 101 |
| Operating in SQL*Plus | 101 |
| Exiting SQL*Plus | 102 |
| SQL*Plus and SQL Commands | 103 |
| SQL*Plus Security..... | 103 |
| Setting the SQL*Plus Environment with the SET Command | 106 |
| Setting Common SQL*Plus Variables | 107 |
| SQL*Plus Error Logging | 111 |
| SQL*Plus Command-Line Options | 113 |
| SQL*Plus Administrative Commands | 115 |
| CLEAR Command | 115 |
| STORE Command | 115 |
| SHOW Command | 116 |
| Key SQL*Plus “Working” Commands | 118 |
| SQLPROMPT Command..... | 118 |
| DESCRIBE Command | 119 |
| HOST Command | 119 |
| SPOOL Command..... | 120 |
| ACCEPT and PROMPT Commands | 121 |
| EXECUTE Command | 121 |
| PAUSE Command | 121 |
| Commands for Formatting Output and Reporting | 122 |
| BREAK Command | 122 |
| COLUMN Command | 123 |
| COMPUTE Command | 123 |
| REPFOOTER Command | 123 |
| REPHEADER Command | 123 |
| BTITLE and TTITLE Commands..... | 123 |

| | |
|--|-----|
| Creating Command Files in SQL*Plus | 124 |
| Saving the SQL Buffer Contents to a File | 124 |
| Executing SQL Scripts in SQL*Plus | 124 |
| Creating a Windows Batch Script | 126 |
| DEFINE and UNDEFINE Commands | 126 |
| Predefined SQL*Plus Variables | 127 |
| Using Comments in SQL*Plus | 128 |
| Listing SQL Commands | 128 |
| Editing Within SQL*Plus | 129 |
| Inserting and Deleting Lines | 130 |
| Adding to Text | 131 |
| Incorporating Comments with the REMARK Command | 132 |
| Copying Tables with the COPY Command | 132 |
| Making DML Changes Permanent with SQL*Plus | 133 |
| Creating Web Pages Using SQL*Plus | 134 |
| Key SQL*Plus Database Administration Commands | 134 |
| RECOVER Command | 134 |
| STARTUP and SHUTDOWN Commands | 134 |
| ARCHIVE LOG Command | 135 |
| Using SQL to Generate SQL | 135 |
| Oracle SQL Developer | 136 |
| Oracle Enterprise Manager | 136 |
| Benefits of Using OEM to Manage Databases | 137 |
| OEM Architecture and Components | 139 |
| Configuring and Using the Database Control | 140 |
| Accessing the Database Control | 143 |
| A Brief Tour of the Database Control | 144 |
| Performance | 145 |
| Availability | 146 |
| Server | 146 |
| Schema | 147 |
| Data Movement | 148 |
| Software and Support | 148 |
| Oracle Software Cloning | 148 |
| Configuring Using the Setup Page | 148 |
| The Related Links Section and the Advisor Central Page | 150 |
| Creating Database Control Roles | 150 |
| Linking to MetaLink | 150 |
| Policy-Based Configuration Framework | 151 |
| Tracking Database Feature-Usage Statistics | 151 |
| OEM Grid Control | 153 |

Managing the Management Agent 156

The Oracle Management Service 157

Connecting to the Grid Control 157

Logging Into the Grid Control..... 157

Features of the Grid Control. 158

Using the Grid Control Home Page 159

Monitoring Your Entire System with the Grid Control 159

PART 2 ■ ■ ■ **Oracle Database 11g Architecture,
Schema, and Transaction
Management**

■ **CHAPTER 5 Oracle Database 11g Architecture 165**

Oracle Database Structures 165

 Logical Database Structures 165

 Physical Database Structures 173

 Other Files 176

Oracle Processes 179

 Interaction Between the User and Oracle Processes 179

 The Server Process..... 179

 The Background Processes..... 180

Oracle Memory Structures 186

 Understanding Main Memory 186

 The System Global Area..... 187

 The Program Global Area..... 193

A Simple Oracle Database Transaction 196

Committing and Rolling Back 197

 Committing a Transaction 197

 Rolling Back a Transaction 198

Data Consistency and Data Concurrency 198

 The Database Writer and the Write Ahead Protocol 199

 The System Change Number..... 199

 Undo Management..... 200

Backup and Recovery Architecture 201

 User-Managed Backup and Recovery 201

 RMAN..... 201

 Oracle Secure Backup 202

 Flashback Recovery Techniques..... 202

| | |
|--|---------|
| The Oracle Data Dictionary and the Dynamic Performance Views | 202 |
| The Oracle Data Dictionary | 203 |
| The Dynamic Performance (V\$) Views | 204 |
| The Three Sets of Data Dictionary Views | 204 |
| How Is the Data Dictionary Created? | 204 |
| Using the Static Data Dictionary Views | 204 |
| The Oracle Optimizer | 205 |
| Talking to the Database | 205 |
| Connecting to Oracle | 205 |
| Oracle Enterprise Manager | 206 |
| SQL*Plus | 207 |
| Oracle Utilities | 207 |
| Data Pump Export and Import | 207 |
| SQL*Loader | 207 |
| LogMiner | 207 |
| Automatic Diagnostic Repository Control Interface | 208 |
| Scheduling and Resource Management | 208 |
| The Oracle Scheduler | 208 |
| Database Resource Manager | 208 |
| Automatic Database Management | 208 |
| Automatic Database Diagnostic Monitor | 209 |
| Automatic Undo Retention Tuning | 209 |
| Automatic Optimizer Statistics Collection | 209 |
| Automatic Storage Management | 209 |
| Automatic SQL Tuning | 209 |
| Common Manageability Infrastructure | 210 |
| Automatic Workload Repository | 210 |
| Active Session History | 210 |
| Server-Generated Alerts | 211 |
| Automated Tasks Feature | 211 |
| Advisory Framework | 212 |
| Change Management | 212 |
| Efficient Managing and Monitoring | 213 |
| CHAPTER 6 Managing Tablespaces | 215 |
| Tablespace Extent Sizing and Space Management | 216 |
| Allocating the Extent Size: Autoallocate vs. Uniform | 216 |
| Automatic vs. Manual Segment Space Management | 217 |

| | |
|--|-----|
| Creating Tablespaces | 218 |
| Data Files and Tablespaces | 219 |
| Extent Allocation and Deallocation | 220 |
| Storage Parameters | 220 |
| Storage Allocation to Database Objects | 222 |
| Adding Space to a Tablespace | 223 |
| Removing Tablespaces | 224 |
| Number of User Tablespaces | 225 |
| Tablespace Quotas | 226 |
| Proactive Tablespace Space Alerts | 226 |
| Managing Logging of Redo Data | 227 |
| Renaming Tablespaces | 228 |
| Read-Only Tablespaces | 229 |
| Taking Tablespaces Offline | 229 |
| Temporary Tablespaces | 229 |
| Creating a Temporary Tablespace | 230 |
| Altering a Temporary Tablespace | 231 |
| Shrinking Temporary Tablespaces | 231 |
| Default Temporary Tablespace | 232 |
| Temporary Tablespace Groups | 233 |
| Default Permanent Tablespaces | 235 |
| Bigfile Tablespaces | 236 |
| Restrictions on Using Bigfile Tablespaces | 237 |
| Creating Bigfile Tablespaces | 237 |
| Altering a Bigfile Tablespace | 238 |
| Viewing Bigfile Tablespace Information | 238 |
| Managing the Sysaux Tablespace | 238 |
| Creating the Sysaux Tablespace | 239 |
| Usage Restrictions for the Sysaux Tablespace | 240 |
| Encrypting Tablespaces | 240 |
| Why You Need Encrypted Tablespaces | 240 |
| Creating the Oracle Wallet | 241 |
| Creating an Encrypted Tablespace | 242 |
| Data Dictionary Views for Managing Tablespaces | 243 |
| DBA_TABLESPACES | 243 |
| DBA_FREE_SPACE | 243 |
| DBA_SEGMENTS | 244 |
| DBA_DATA_FILES | 245 |
| DBA_TABLESPACE_GROUPS | 246 |
| V\$DATAFILE | 246 |
| V\$FILESTAT | 246 |

| | |
|--|-----|
| Easy File Management with Oracle Managed Files | 247 |
| Benefits of Using OMF | 248 |
| Creating Oracle Managed Files | 248 |
| Different Types of Oracle Managed Files | 250 |
| Copying Files Between Two Databases | 253 |
| COPY_FILE | 253 |
| GET_FILE | 254 |
| PUT_FILE | 254 |
| Finding Out How Much Free Space Is Left | 255 |
| Working with Operating System Files | 256 |
| Using the UTL_FILE Package | 256 |
| Key UTL_FILE Procedures and Functions | 257 |
| Exceptions | 258 |
| A Simple Example Using the UTL_FILE Package | 258 |

CHAPTER 7 **Schema Management** 261

| | |
|---|-----|
| Types of SQL Statements | 261 |
| System-Control Statements | 262 |
| Session-Control Statements | 262 |
| Embedded SQL Statements | 262 |
| Data Manipulation Language Statements | 263 |
| Transaction-Control Statements | 263 |
| Data Definition Language Statements | 264 |
| Oracle Schemas | 264 |
| Oracle Tables | 265 |
| Estimating the Table Size | 266 |
| Creating a Simple Table | 268 |
| What Is a Null Value? | 269 |
| Default Values for Columns | 270 |
| Virtual Columns | 270 |
| Adding a Column to a Table | 271 |
| Dropping a Column from a Table | 271 |
| Renaming a Table Column | 272 |
| Renaming a Table | 272 |
| Removing All Data from a Table | 272 |
| Creating a New Table with the CTAS Option | 273 |
| Placing a Table in Read-Only Mode | 273 |
| Table Compression | 274 |
| Dropping Tables | 276 |

| | |
|---|-----|
| Special Oracle Tables | 277 |
| Temporary Tables | 277 |
| Index-Organized Tables | 278 |
| External Tables | 280 |
| Partitioned Tables | 280 |
| Range Partitioning | 281 |
| Interval Partitioning | 282 |
| Hash Partitioning | 283 |
| List Partitioning | 284 |
| Reference Partitioning | 284 |
| Virtual Column-Based Partitioning | 286 |
| System Partitioning | 287 |
| Composite Partitioning | 287 |
| Partition Maintenance Operations | 290 |
| Data Dictionary Views for Managing Tables | 292 |
| Clusters | 295 |
| Hash Clusters | 296 |
| Oracle Indexes | 296 |
| Guidelines for Creating Indexes | 297 |
| Oracle Index Schemes | 298 |
| Estimating the Size of an Index | 298 |
| Creating an Index | 299 |
| Special Types of Indexes | 300 |
| Bitmap Indexes | 301 |
| Reverse-Key Indexes | 301 |
| Key-Compressed Indexes | 301 |
| Function-Based Indexes | 302 |
| Partitioned Indexes | 302 |
| Invisible Indexes | 303 |
| Monitoring Index Usage | 304 |
| Index Maintenance | 305 |
| Managing Database Integrity Constraints | 306 |
| Primary Key Constraints | 306 |
| Not Null Constraints | 307 |
| Check Constraints | 307 |
| Unique Constraints | 307 |
| Referential Integrity Constraints | 308 |
| Integrity Constraint States | 308 |
| Rely Constraints | 310 |
| Deferrable and Immediate Constraints | 310 |

| | |
|---|-----|
| Constraint- and Index-Related Views | 310 |
| DBA_CONSTRAINTS | 310 |
| DBA_CONS_COLUMNS | 311 |
| Using Views | 312 |
| Using Materialized Views | 314 |
| Query Rewriting | 315 |
| The Rewrite_or_Error Hint | 315 |
| Rewrite Integrity | 316 |
| Refreshing Materialized View Data | 316 |
| Using the DBMS_MVIEW Package | 317 |
| Creating Materialized Views | 317 |
| Creating the Materialized View | 318 |
| Using the SQL Access Advisor | 320 |
| Using the OEM Database Control | 320 |
| Using the DBMS_ADVISOR Package | 323 |
| Using the QUICK_TUNE Procedure | 324 |
| Using Synonyms | 324 |
| Creating a Public Synonym | 325 |
| Creating a Private Synonym | 326 |
| Dropping a Synonym | 326 |
| Managing Synonyms | 326 |
| Switching to a Different Schema | 327 |
| Using Sequences | 327 |
| Using Triggers | 328 |
| Viewing Object Information | 329 |
| Views for Managing Database Objects | 329 |
| DBA_OBJECTS | 329 |
| DBA_TABLES | 330 |
| DBA_EXTERNAL_TABLES | 330 |
| DBA_TAB_PARTITIONS | 330 |
| DBA_PART_TABLES | 331 |
| DBA_TAB_MODIFICATIONS | 331 |
| DBA_TAB_COLUMNS | 332 |
| DBA_VIEWS | 332 |
| DBA_MVIEWS | 333 |
| DBA_INDEXES | 333 |
| DBA_IND_COLUMNS | 333 |
| INDEX_STATS | 334 |

| | | |
|------------------|---|------------|
| CHAPTER 8 | Oracle Transaction Management | 337 |
| | Oracle Transactions | 337 |
| | COMMIT Statement | 338 |
| | ROLLBACK Statement | 339 |
| | Transaction Properties | 340 |
| | Transaction Concurrency Control | 341 |
| | Concurrency Problems | 341 |
| | Schedules and Serializability | 342 |
| | Isolation Levels and the ISO Transaction Standard | 342 |
| | Oracle's Isolation Levels | 344 |
| | Transaction- and Statement-Level Consistency | 345 |
| | Changing the Default Isolation Level | 345 |
| | Implementing Oracle's Concurrency Control | 347 |
| | Oracle Locking Methods | 348 |
| | Oracle Lock Types | 348 |
| | Allowing DDL Locks to Wait for DML Locks | 350 |
| | Explicit Table Locking | 350 |
| | Explicit Locking in Oracle | 351 |
| | Managing Oracle Locks | 353 |
| | Using Undo Data to Provide Read Consistency | 356 |
| | Automatic Undo Management | 356 |
| | Guaranteed Undo Retention | 362 |
| | Using the OEM to Manage Undo Data | 365 |
| | Flashback Error Correction Using Undo Data | 366 |
| | Querying Old Data with Flashback Query | 367 |
| | Flashback Using the DBMS_FLASHBACK Package | 368 |
| | Flashback Versions Query | 369 |
| | Flashback Transaction Query | 372 |
| | Using Flashback Transaction Query and Flashback | |
| | Versions Query Together | 375 |
| | The Flashback Table Feature | 376 |
| | Transaction Management | 380 |
| | Discrete Transactions | 380 |
| | Autonomous Transactions | 380 |
| | Resumable Space Allocation | 382 |
| | Resumable Operations | 383 |
| | Common Resumable Errors | 383 |
| | Using the Resumable Space Allocation Feature | 383 |
| | Notification of Suspended Operations | 386 |
| | Operation-Suspended Alert | 386 |
| | Monitoring Resumable Space Allocation | 386 |

| | |
|---|-----|
| Managing Long Transactions | 386 |
| Benefits of Using the Workspace Manager | 387 |
| Table Versioning and Workspaces | 387 |

PART 3 ■■■ Installing Oracle Database 11g, Upgrading, and Creating Databases

| | | |
|-------------|--|------------|
| ■ CHAPTER 9 | Installing and Upgrading to Oracle Database 11g | 391 |
| | Installing Oracle | 391 |
| | Reviewing the Documentation | 392 |
| | Determining Disk and Memory Requirements..... | 392 |
| | Optimal Flexible Architecture | 393 |
| | Mount Points..... | 394 |
| | Directory and File-Naming Conventions..... | 394 |
| | Performing Preinstallation Tasks | 400 |
| | Checking the Preinstallation Requirements..... | 400 |
| | System Administrator's Preinstallation Tasks..... | 401 |
| | Oracle Owner's Preinstallation Tasks..... | 410 |
| | A Final Checklist for the Installation | 413 |
| | Accessing the Installation Software | 414 |
| | Installing the Software | 416 |
| | Using Response Files to Install Oracle Software | 421 |
| | After the Installation | 422 |
| | System Administrator's Post-Installation Tasks | 423 |
| | Oracle Owner's Post-Installation Tasks | 424 |
| | Uninstalling Oracle | 425 |
| | Removing All Oracle Databases Running on Your Server | 425 |
| | Removing the Oracle Software | 426 |
| | Upgrading to Oracle Database 11g | 426 |
| | Routes to Oracle Database 11g | 426 |
| | Upgrade Methods and Tools | 427 |
| | The Manual Upgrade Process | 427 |
| | The Database Upgrade Assistant | 428 |
| | The Pre-Upgrade Information Tool | 428 |
| | The Post-Upgrade Status Tool..... | 429 |
| | Preparing the Database for the Upgrade | 430 |

| | |
|--|-----|
| Upgrading with the DBUA | 430 |
| Starting the DBUA | 430 |
| Running the DBUA | 431 |
| Restoring the Pre-Upgrade Database | 433 |
| Upgrading Manually | 434 |
| Upgrade and Downgrade Scripts | 434 |
| Creating a Spool File | 434 |
| Running the Pre-Upgrade Information Tool | 435 |
| Backing up the Database | 437 |
| Copying the Parameter File | 437 |
| Starting Up the New Database | 437 |
| Running the Upgrade Script | 438 |
| Running the Upgrade Actions Script | 438 |
| Restarting the Instance | 438 |
| Running the Post Upgrade Actions Script | 439 |
| Checking for Invalid Objects | 439 |
| Recompiling and Validating Objects | 439 |
| Running the Post-Upgrade Status Tool | 440 |
| Ending the Spool File | 441 |
| Restarting the New Database | 441 |
| After the Upgrade | 441 |
| Resetting Passwords for Stronger Password Protection | 441 |
| Downgrading to an Old Version | 441 |

CHAPTER 10 Creating a Database 443

| | |
|---|-----|
| Getting Ready to Create the Database | 443 |
| Installing the Oracle Software | 444 |
| Creating the File System for the Database | 444 |
| Ensuring Sufficient Memory Allocation | 445 |
| Getting Necessary Authorizations | 446 |
| Setting the OS Environment Variables | 446 |
| Creating the Parameter File | 446 |
| Types of Database Parameter Files | 446 |
| The Initialization Parameter File | 447 |
| Changing the Initialization Parameter Values | 448 |
| Important Oracle Database 11g Initialization Parameters | 449 |
| Undocumented Initialization Parameters | 473 |
| Viewing the Current Initialization Parameter Values | 473 |

| | |
|--|-----|
| Creating a New Database | 474 |
| Manual Creation | 474 |
| Using the DBCA to Create a Database | 486 |
| Using a Server Parameter File | 493 |
| Creating a Server Parameter File | 494 |
| Setting the Scope of Dynamic Parameter Changes | 496 |
| Creating an SPFILE or PFILE from Memory | 497 |
| Starting Up and Shutting Down the Database from SQL*Plus | 497 |
| Starting the Database..... | 497 |
| Automatically Starting Databases..... | 499 |
| Restricting Database Access..... | 501 |
| Shutting Down the Database..... | 502 |
| Quiescing a Database..... | 505 |
| Suspending a Database..... | 505 |
| Dropping a Database | 506 |
| Using the Data Dictionary to Monitor Database Status..... | 507 |
| What Next? | 507 |

PART 4 ■■■ **Connectivity and User Management**

■ **CHAPTER 11 Connectivity and Networking** 511

| | |
|--|-----|
| Oracle Networking and Database Connectivity | 511 |
| Networking Concepts: How Oracle Networking Works | 513 |
| How a Web Application Connects to Oracle..... | 513 |
| Database Instance Names..... | 514 |
| Global Database Names..... | 514 |
| Database Service Names..... | 514 |
| Connect Descriptors | 514 |
| Connect Identifiers | 515 |
| Connect Strings..... | 515 |
| Establishing Oracle Connectivity | 516 |
| The Oracle Client | 517 |
| Installing the Oracle Client..... | 518 |
| Using the TWO_TASK Environment Variable..... | 519 |
| The Instant Client | 519 |
| Installing the Instant Client | 519 |

- The Listener and Connectivity 520
 - Automatic Service Registration 521
 - Listener Commands 522
 - Listener Management 523
- Naming and Connectivity 525
 - The Local Naming Method 525
 - The Easy Connect Naming Method 529
 - Database Resident Connection Pooling 531
 - The External Naming Method 533
 - The Directory Naming Method 534
- Oracle and Java Database Connectivity 537
 - Establishing Database Connectivity 538
 - Working with the Database 538
 - A Complete Program 541

CHAPTER 12 User Management and Database Security 543

- Managing Users 544
 - Temporary and Default Tablespaces 544
 - Creating a New User 544
 - Altering a User 547
 - Dropping a User 547
 - Creating and Using User Profiles 548
 - Managing Resources 554
- The Database Resource Manager 554
 - Using the Database Resource Manager 555
 - Using OEM to Administer the Database Resource Manager 566
- Controlling Database Access 567
 - Privileges in an Oracle Database 567
 - Roles 574
 - Using Views and Stored Procedures to Manage Privileges 577
 - DBA Views for Managing Users, Roles, and Privileges 577
 - Fine-Grained Data Access 578
- Auditing Database Usage 586
 - Standard Auditing 587
 - Fine-Grained Auditing 593
- Authenticating Users 596
 - Database Authentication 596
 - External Authentication 601
 - Proxy Authentication 602
 - Centralized User Authorization 602

| | |
|--|-----|
| Enterprise User Security | 603 |
| Shared Schemas | 603 |
| Single Sign-On | 603 |
| Data Encryption | 603 |
| Tablespace Encryption | 608 |
| Oracle Internet Directory | 611 |
| Database Security Dos and Don'ts | 611 |
| Automatic Secure Configuration | 611 |
| User Accounts | 611 |
| Passwords | 611 |
| Operating System Authentication | 612 |
| Database Auditing | 612 |
| Granting Privileges | 612 |
| Dealing with Environments with Multiple DBAs | 613 |
| Protecting the Data Dictionary | 613 |
| Setting Permissions | 613 |
| The Network and the Listener | 614 |
| Fine-Grained Network Access Control | 615 |
| Oracle's Advanced Security Option | 618 |
| Application Security | 618 |
| Useful Techniques for Managing Users | 619 |

PART 5 ■■■ **Data Loading, Backup, and Recovery**

| | |
|--|-----|
| CHAPTER 13 Loading and Transforming Data | 625 |
| An Overview of Extraction, Transformation, and Loading | 625 |
| Using the SQL*Loader Utility | 627 |
| Exploring the SQL*Loader Control File | 628 |
| Generating Data During the Load | 636 |
| Invoking SQL*Loader | 637 |
| Exploring the Loader Log File | 638 |
| Using Return Codes | 639 |
| Using the Direct-Path Loading Method | 639 |
| Some Useful SQL*Loader Data-Loading Techniques | 642 |

| | |
|--|-----|
| Using External Tables to Load Data | 645 |
| Creating the External Table Layer..... | 646 |
| Populating External Tables | 649 |
| Using an External Table | 652 |
| Using SQL*Loader to Generate External Table Creation Statements | 653 |
| Transforming Data | 656 |
| Deriving the Data from Existing Tables | 656 |
| Using SQL to Transform Data | 657 |
| Using the SQL MODEL Clause..... | 666 |
| Using Oracle Streams for Replication and Information Sharing | 670 |
| Exploring the Streams Architecture | 671 |
| Setting Up Oracle Streams..... | 671 |

CHAPTER 14 Using Data Pump Export and Import **677**

| | |
|--|-----|
| Introduction to the Data Pump Technology | 677 |
| Benefits of the Data Pump Technology | 678 |
| Uses for Data Pump Export and Import | 679 |
| Data Pump Components | 680 |
| Data-Access Methods | 680 |
| Data Pump Files | 681 |
| Data Pump Privileges | 685 |
| The Mechanics of a Data Pump Job..... | 685 |
| Performing Data Pump Exports and Imports | 686 |
| Data Pump Export Methods..... | 687 |
| Data Pump Export Modes | 688 |
| Data Pump Export Parameters | 689 |
| Data Pump Export Examples..... | 704 |
| Data Pump Import Types and Modes | 705 |
| Data Pump Import Parameters | 705 |
| Monitoring a Data Pump Job | 713 |
| Viewing Data Pump Jobs..... | 713 |
| Viewing Data Pump Sessions | 714 |
| Viewing Data Pump Job Progress..... | 714 |
| Using the Data Pump API | 715 |
| Transportable Tablespaces | 716 |
| Uses for Transportable Tablespaces | 716 |
| Transporting a Tablespace | 716 |

| | |
|--|------------|
| Transporting Tablespaces Across Platforms with Different Endian Formats | 720 |
| CHAPTER 15 Backing Up Databases | 725 |
| Backing Up Oracle Databases | 725 |
| Important Backup Terms | 726 |
| Backup Guidelines | 728 |
| Testing Your Backups | 730 |
| Maintaining a Redundancy Set | 730 |
| Backup Strategies | 731 |
| The Flash Recovery Area | 734 |
| Benefits of the Flash Recovery Area | 734 |
| Setting the Size of the Flash Recovery Area | 736 |
| Ways to Create a Flash Recovery Area | 736 |
| Backing Up the Flash Recovery Area | 739 |
| Working with the Flash Recovery Area | 740 |
| Recovery Manager | 741 |
| Benefits of RMAN | 742 |
| RMAN Architecture | 743 |
| Connecting to RMAN | 745 |
| Scripting with RMAN | 747 |
| Important RMAN Terms | 752 |
| Specifying Backup Tags and Backup Formats | 753 |
| Making Copies of RMAN Backups | 754 |
| RMAN Backup Locations | 755 |
| RMAN Commands | 755 |
| RMAN Configuration Parameters | 761 |
| Working with the Recovery Catalog | 766 |
| Maintaining the Recovery Catalog | 769 |
| Virtual Private Catalogs | 772 |
| Examples of RMAN Backups | 774 |
| Restarting an RMAN Backup | 777 |
| Specifying Limits for Backup Duration | 777 |
| Incrementally Updated Backups | 778 |
| Fast Incremental Backups | 779 |
| RMAN Compressed Backups | 780 |
| Archival Backups | 780 |
| Monitoring and Verifying RMAN Jobs | 782 |
| Backing Up the Control File | 784 |

| | |
|--|-----|
| Oracle Secure Backup | 785 |
| Benefits of Oracle Secure Backup | 786 |
| Oracle Secure Backup Administrative Domain | 786 |
| Installing Oracle Secure Backup | 786 |
| Using the Oracle Backup Web Interface Tool | 788 |
| Configuring Oracle Secure Backup | 788 |
| Performing Backups with Oracle Secure Backup | 789 |
| User-Managed Backups | 790 |
| Making Whole Database Backups | 790 |
| Making Partial Database Backups | 794 |
| Monitoring User-Managed Online Backups | 794 |
| Database Corruption Detection | 795 |
| Detecting Media Corruption | 795 |
| Detecting Data Block Corruption | 795 |
| Enhanced Data Protection for Disaster Recovery | 798 |
| High-Availability Systems | 798 |
| Oracle Data Guard and Standby Databases | 799 |

CHAPTER 16 Database Recovery

| | |
|--|-----|
| Types of Database Failures | 801 |
| System Failures | 802 |
| Data Center Disasters | 802 |
| Human Error | 802 |
| Media Failures | 802 |
| Failures and Data Repair | 803 |
| The Oracle Recovery Process | 804 |
| Crash and Instance Recovery | 804 |
| Faster Instance Startup | 805 |
| Media Recovery | 806 |
| Media Recovery vs. Nonmedia Recoveries | 808 |
| Performing Recovery with RMAN | 809 |
| RMAN's Advantages for Recovery | 810 |
| Using VALIDATE BACKUP to Validate RMAN Backups | 810 |
| Using the RESTORE . . . VALIDATE Command | 811 |
| Using the RESTORE . . . PREVIEW Command | 811 |
| Identifying Necessary Files for Recovery | 812 |
| RMAN Recovery Procedures | 812 |
| Recovering with Incrementally Updated Backups | 812 |

| | |
|---|-----|
| Monitoring RMAN Jobs | 813 |
| User-Managed Recovery Procedures | 813 |
| Typical Media Recovery Scenarios | 814 |
| Complete Recovery of a Whole Database | 814 |
| Recovering a Tablespace | 817 |
| Recovering a Datafile | 818 |
| Incomplete Recovery | 820 |
| Recovering from the Loss of Control Files | 824 |
| Recovering a Datafile Without a Backup | 828 |
| The Data Recovery Advisor | 829 |
| Working with the Data Recovery Advisor | 830 |
| Cloning a Database | 833 |
| Using RMAN to Clone a Database | 834 |
| Using Database Control to Clone a Database | 838 |
| Manually Cloning a Database | 839 |
| Techniques for Granular Recovery | 840 |
| Tablespace Point-in-Time Recovery | 840 |
| Using LogMiner for Precision Recovery | 841 |
| Flashback Techniques and Recovery | 847 |
| Flashback Levels | 848 |
| Flashback vs. Traditional Recovery Techniques | 848 |
| Flashback Drop | 849 |
| Flashback Database | 853 |
| Using Restore Points | 861 |
| Guaranteed Restore Points | 862 |
| Viewing Restore Points | 863 |
| Repairing Data Corruption and Trial Recovery | 864 |
| Block Media Recovery | 864 |
| Trial Recovery | 865 |
| Troubleshooting Recovery Errors | 866 |
| The ORA-01194 Error | 866 |
| The ORA-01152 Error | 867 |
| The ORA-00376 Error | 868 |
| The Transaction Backout Feature | 868 |
| Flashback Data Archive | 870 |
| Managing the Flashback Data Archive | 870 |
| Enabling the Flashback Data Archive | 872 |
| Using Flashback Data Archives: Examples | 873 |

PART 6 ■ ■ ■ Managing the Database

■ CHAPTER 17 Automatic Management and Online Capabilities 877

 The Automatic Database Diagnostic Monitor 877

 The Purpose of the ADDM 878

 Problems That the ADDM Diagnoses 878

 ADDM Findings 880

 ADDM Recommendations 880

 Managing the ADDM 881

 Automatic Memory Management 894

 Different Types of Memory Management 894

 Enabling Automatic Memory Management 895

 Tuning Automatic Memory Management 896

 Automatic Memory Parameter Dependency 897

 Automatic Optimizer Statistics Collection 897

 Automatically Collecting Optimizer Statistics 898

 Manually Collecting Optimizer Statistics 900

 Automatic Storage Management 900

 Benefits of ASM 901

 Examining the ASM Architecture 901

 Installing ASM 902

 ASM and Cluster Synchronization Service 902

 Creating an ASM Instance 904

 The ASM Fast Mirror Resync Feature 908

 ASM Preferred Mirror Read 909

 Changing ASM Disk Group Attributes 909

 The asmcmd Command-Line Tool 911

 Managing ASM Disk Groups 913

 Adding Performance and Redundancy with Disk Groups 913

 Creating a Disk Group 914

 Adding Disks to a Disk Group 915

 Dropping Disks and Disk Groups 916

 Rebalancing Disk Groups 916

 Managing ASM Files 916

 Types of ASM Filenames 917

 Working with ASM Filenames 918

 ASM File Templates 918

 Creating an ASM-Based Database 919

 Migrating Your Database to ASM 919

| | |
|---|-----|
| Automatic Space Management | 921 |
| Automatic Undo Management..... | 921 |
| Easy File Management with OMF..... | 922 |
| Online Segment Shrinking and the Segment Advisor..... | 927 |
| Manual Segment Shrinking | 928 |
| Using the Segment Advisor to Shrink Segments..... | 930 |
| Automatic Segment Advisor Job..... | 931 |
| Automatic Checkpoint Tuning | 932 |
| Online Capabilities of Oracle Database 11g..... | 933 |
| Online Data Reorganization | 933 |
| Online Data Redefinition | 935 |
| Dynamic Resource Management | 941 |
| Online Database Block-Size Changes | 943 |
| Using Database Quiescing for Online Maintenance | 944 |
| Suspending the Database | 945 |

■ CHAPTER 18 Managing and Monitoring the Operational Database ... 947

| | |
|---|-----|
| Types of Oracle Performance Statistics | 948 |
| Cumulative Statistics | 948 |
| Sample Data | 948 |
| Baseline Data | 948 |
| Database Metrics | 950 |
| Server-Generated Alerts | 952 |
| Default Server-Generated Alerts..... | 952 |
| Baseline Metrics and Adaptive Thresholds | 953 |
| Managing Alerts | 954 |
| Proactive Tablespace Alerts | 956 |
| Using the Alert Logs and Trace Files for Monitoring..... | 958 |
| Data Dictionary Views Related to Metrics and Alerts | 958 |
| The Automatic Workload Repository | 959 |
| Types of Data Collected by the AWR | 960 |
| AWR Data Handling | 960 |
| Managing the AWR..... | 961 |
| Moving Window Baselines..... | 965 |
| AWR Baseline Templates..... | 965 |
| Active Session History | 971 |
| Current Active Session Data | 972 |
| Older Active Session History Data..... | 972 |
| Producing an ASH Report..... | 972 |

| | |
|---|------|
| The Management Advisory Framework | 975 |
| The Management Advisors | 976 |
| Managing the Advisory Framework | 977 |
| Working with the Undo and the MTTR Advisors | 980 |
| Using the Undo Advisor | 980 |
| Using the MTTR Advisor | 981 |
| Managing Online Redo Logs | 981 |
| Hardware Mirroring vs. Oracle Multiplexing | 982 |
| Online Redo Log Groups | 982 |
| Creating Online Redo Log Groups | 982 |
| Adding Redo Log Groups | 983 |
| Renaming Redo Log Files | 983 |
| Dropping Online Redo Logs | 984 |
| Online Redo Log Corruption | 984 |
| Monitoring the Redo Logs | 984 |
| Managing Database Links | 985 |
| Creating a Private Database Link | 985 |
| Creating a Public Database Link | 986 |
| Using the Database Control to Create Database Links | 987 |
| Comparing and Converging Database Objects | 987 |
| Comparing Data | 987 |
| Converging Data | 990 |
| Copying Files with the Database Server | 991 |
| Requirements for the File Copy | 991 |
| Copying Files on a Local System | 991 |
| Transferring a File to a Different Database | 992 |
| Mapping Oracle Files to Physical Devices | 993 |
| Architecture of File Mapping | 993 |
| Setting Up File Mapping | 993 |
| Using the Oracle Scheduler | 994 |
| Basic Scheduler Components | 994 |
| Types of Scheduler Jobs | 995 |
| Advanced Scheduler Components | 996 |
| Scheduler Architecture | 997 |
| Scheduler Privileges | 997 |
| Managing the Basic Scheduler Components | 998 |
| Managing Lightweight Jobs | 1000 |
| Managing External Jobs | 1002 |
| Managing Advanced Scheduler Components | 1011 |
| Managing Scheduler Attributes | 1017 |

| | |
|--|------|
| Purging Job Logs | 1019 |
| Default Scheduler Jobs | 1019 |
| Automated Maintenance Tasks | 1019 |
| Predefined Maintenance Windows | 1020 |
| Managing Automated Maintenance Tasks | 1020 |
| Monitoring Automated Maintenance Tasks | 1020 |
| Fault Diagnosability | 1022 |
| Automatic Diagnostic Repository | 1023 |
| ADRCI | 1024 |
| Incident Packaging Service | 1026 |
| The Support Workbench | 1028 |
| The Health Monitor | 1032 |
| Repairing SQL Statements with the SQL Repair Advisor | 1035 |
| The SQL Test Case Builder | 1038 |

PART 7 ■■■ Performance Tuning

■ CHAPTER 19 Improving Database Performance:

| | |
|---|------|
| SQL Query Optimization | 1041 |
| An Approach to Oracle Performance Tuning | 1041 |
| A Systematic Approach to Performance Tuning | 1042 |
| Reactive Performance Tuning | 1042 |
| Optimizing Oracle Query Processing | 1043 |
| Parsing | 1043 |
| Optimization | 1043 |
| Query Execution | 1046 |
| Query Optimization and the Oracle CBO | 1047 |
| Choosing Your Optimization Mode | 1047 |
| Providing Statistics to the Optimizer | 1047 |
| Setting the Optimizer Mode | 1049 |
| Setting the Optimizer Level | 1050 |
| What Does the Optimizer Do? | 1051 |
| Drawbacks of the CBO | 1053 |
| Providing Statistics to the CBO | 1053 |
| Deferred Statistics Publishing | 1056 |
| Extended Statistics | 1058 |
| The Cost Model of the Oracle Optimizer | 1060 |
| Collecting System Statistics | 1060 |

| | |
|---|------|
| Collecting Statistics on Dictionary Objects | 1062 |
| Frequency of Statistics Collection | 1063 |
| What Happens When You Don't Have Statistics | 1063 |
| Using the OEM to Collect Optimizer Statistics | 1064 |
| Writing Efficient SQL | 1065 |
| Efficient WHERE Clauses | 1065 |
| Using Hints to Influence the Execution Plan | 1067 |
| Selecting the Best Join Method | 1068 |
| Using Bitmap Join Indexes | 1069 |
| Selecting the Best Join Order | 1070 |
| Indexing Strategy | 1070 |
| Monitoring Index Usage | 1073 |
| Removing Unnecessary Indexes | 1073 |
| Using Similar SQL Statements | 1074 |
| Reducing SQL Overhead Via Inline Functions | 1074 |
| Using Bind Variables | 1075 |
| Avoiding Improper Use of Views | 1075 |
| Avoiding Unnecessary Full Table Scans | 1075 |
| How the DBA Can Help Improve SQL Processing | 1075 |
| Using Partitioned Tables | 1076 |
| Using Compression Techniques | 1076 |
| Using Materialized Views | 1077 |
| Using Stored Outlines to Stabilize the CBO | 1077 |
| SQL Plan Management | 1080 |
| SQL Plan Baselines | 1080 |
| Capturing SQL Plan Baselines | 1081 |
| The SQL Management Base | 1085 |
| Using Parallel Execution | 1085 |
| Other DBA Tasks | 1086 |
| Adaptive Cursor Sharing | 1087 |
| How Adaptive Cursor Sharing Works | 1088 |
| Monitoring Adaptive Cursor Sharing | 1089 |
| SQL Performance Tuning Tools | 1090 |
| Using EXPLAIN PLAN | 1090 |
| Using Autotrace | 1095 |
| Using SQL Trace and TKPROF | 1099 |
| End-to-End Tracing | 1105 |
| Using the DBMS_MONITOR Package | 1106 |
| Using the V\$SQL View to Find Inefficient SQL | 1108 |

| | |
|--|------|
| The SQL Tuning Advisor | 1111 |
| Using the SQL Tuning Advisor | 1111 |
| How the SQL Tuning Advisor Works | 1111 |
| The SQL Tuning Advisor in Practice | 1113 |
| The Automatic SQL Tuning Advisor | 1115 |
| Interpreting Automatic SQL Tuning Reports | 1119 |
| Using Other GUI Tools | 1120 |
| Using the Result Cache | 1120 |
| Managing the Result Cache | 1120 |
| Setting the RESULT_CACHE_MODE Parameter | 1121 |
| Managing the Result Cache | 1123 |
| The PL/SQL Function Result Cache | 1124 |
| The Client Query Result Cache | 1125 |
| A Simple Approach to Tuning SQL Statements | 1126 |
| Identify Problem Statements | 1127 |
| Locate the Source of the Inefficiency | 1127 |
| Tune the Statement | 1127 |
| Compare Performance | 1127 |


CHAPTER 20 Performance Tuning: Tuning the Instance 1129

| | |
|--|------|
| An Introduction to Instance Tuning | 1129 |
| Automatic Performance Tuning vs. Dynamic Performance Views ... | 1131 |
| Tuning Oracle Memory | 1132 |
| Tuning the Shared Pool | 1133 |
| Hard Parsing and Soft Parsing | 1135 |
| Tuning the Buffer Cache | 1144 |
| Tuning the Large Pool, Streams Pool, and Java Pool | 1148 |
| Tuning PGA Memory | 1148 |
| Evaluating System Performance | 1152 |
| CPU Performance | 1153 |
| Memory | 1158 |
| Disk I/O | 1158 |
| Measuring I/O Performance | 1159 |
| Is the I/O Optimally Distributed? | 1159 |
| Reducing Disk Contention | 1160 |
| The Oracle SAME Guidelines for Optimal Disk Usage | 1160 |
| Network Performance | 1160 |

| | |
|--|------|
| Measuring Instance Performance | 1161 |
| Database Hit Ratios | 1161 |
| Database Wait Statistics | 1162 |
| Wait Events and Wait Classes | 1163 |
| Analyzing Instance Performance | 1164 |
| Examining System Performance | 1181 |
| Know Your Application | 1182 |
| Using the ADDM to Analyze Performance Problems | 1183 |
| Using AWR Reports for Individual SQL Statements | 1184 |
| Operating System Memory Management | 1186 |
| Analyzing Recent Session Activity with an ASH Report | 1186 |
| When a Database Hangs | 1186 |
| Handling a Stuck Archiver Process | 1187 |
| System Usage Problems | 1188 |
| Excessive Contention for Resources | 1188 |
| Locking Issues | 1189 |
| Abnormal Increase in Process Size | 1190 |
| Delays Due to Shared Pool Problems | 1191 |
| Problems Due to Bad Statistics | 1191 |
| Collecting Information During a Database Hang | 1191 |
| Using the Database Control's Hang Analysis Page | 1192 |
| A Simple Approach to Instance Tuning | 1194 |
| What's Happening in the Database? | 1194 |
| Using the OEM Database Control to Examine | |
| Database Performance | 1195 |
| Are There Any Long-Running Transactions? | 1202 |
| Is Oracle the Problem? | 1203 |
| Is the Network Okay? | 1203 |
| Is the System CPU Bound? | 1203 |
| Is the System I/O Bound? | 1204 |
| Is the Database Load Too High? | 1205 |
| Checking Memory-Related Issues | 1205 |
| Are the Redo Logs Sized Correctly? | 1205 |
| Is the System Wait Bound? | 1206 |
| The Compare Periods Report | 1206 |
| Eliminating the Contention | 1208 |
| Real Application Testing | 1209 |
| Database Replay | 1209 |
| SQL Performance Analyzer | 1216 |

| | | |
|-----------------|--|-------------|
| APPENDIX | Oracle Database 11<i>g</i> SQL and PL/SQL: A Brief Primer | 1221 |
| INDEX | | 1255 |

About the Author



SAM ALAPATI is currently an independent Oracle consultant, working at ERCOT (Electric Reliability Council of Texas). Prior to this, Sam was an Oracle DBA at the Boy Scouts of America, Sabre, Oracle Corporation, NBC, and AT&T. Sam has previously written *Expert Oracle9i Database Administration* and *Expert Oracle Database 10g Administration*, both for Apress, as well as three OCP certification books for Oracle Press. Sam has also coauthored *Oracle Database 11g RMAN Recipes* and *Oracle Database 11g New Features for DBAs and Developers*, both published by Apress. Sam has been working with relational databases since 1985, starting with the Ingres database.

About the Technical Reviewer

■ **JOHN WATSON** has spent 20 years hacking away at the Oracle product set, some of those years working for Oracle Corporation (in Internal Support and for Oracle University). Now he works for BPLC Management Consultants, consulting and teaching Oracle server technologies throughout Europe and Africa. He has written and edited several books for Apress and Oracle Press. He currently lives in Oxfordshire with his wife, two Chow Chows, six cats, and a parrot.

Acknowledgments

A book is never the product of a single-handed effort by the author. All books, especially technical books, are the product of intensive effort and care by a number of people, and the author is invariably the beneficiary of the hard work and help of a team of people who labor behind the scenes. So, let me start from the beginning and acknowledge my debts to all those who contributed to making this book a reality. First, I must acknowledge the book's Lead Editor, Jonathan Gennick, well known to many folks in the Oracle field, for explaining complex new topics to Oracle DBAs and developers in several articles and best-selling books. Jonathan is unique among editors: he always makes the author feel like his or her book is the only one he is working with, whereas the truth is that at any given time, Jonathan is probably grappling with a number of books at various stages of writing. Not only has Jonathan given freely of his time and counsel over the last year, but his unfailing courtesy and concern for this author has helped quite a bit. I am forever in the debt of the marvelous technical editing performed by John Watson, who was the technical reviewer for the book. John, as he has done before for some of my other books, has saved me from committing several errors—not only has he pointed out errors, but in most cases, he also pointed out what might work. John's exceptional consciousness and his painstaking testing of various features has truly enhanced the quality of the book. Any errors that have eluded John's keen eye are, alas, my own, and I must take responsibility for them.

This is a long book by any standards, and big books test the patience and the stamina of the best of the people who work on them. I must say that the Apress team came out with flying colors, handling potential problems with aplomb and moving the production process along through a long and arduous period lasting about a year. Project Manager Richard Dal Porto made sure none of us dropped the ball, by gently coaxing us to deliver the chapters on schedule. Richard, thanks for your patience and the skill with which you have led us over this long course. Douglas Pundick, Developmental Editor, had to labor through the extra long chapters, a task he performed with great distinction. I've benefited from Douglas's suggestions regarding chapter styling and ordering of topics. Copy Editor Ami Knox is in a class all by herself. I've had the fortune over the years to work with several good copy editors, but Ami is by far the best. The truly great copy editors don't merely catch errors in style and substance and typos—as Ami has shown me for the past several months, they also have an innate sense of what the author is trying to say and whether the author has managed to convey that message to the readers. Ami's brilliant copyediting has helped me produce a final book that is more readable and, hopefully, contains few errors. I've worked with Production Editor Laura Cheu before, and her crisp professionalism and meticulousness has always enhanced my books and I'm thankful to her. April Eddy, the proofreader, has, as usual, performed a sterling job, catching some hard-to-find potential errors, and I thank her for her hard work. Thanks also to the rest of the production team as well as the marketing team for producing this book.

Moving on to friends and colleagues, I gratefully acknowledge the support of David Campbell, Technical Services Director at Boy Scouts of America (where I've worked until recently for many years). Dave's confidence in me and his helpful nature has always provided an incentive for me to do the best job I could at work and on my books. David Jeffress has been a wonderful manager and a valuable friend over the years and I appreciate all the help David gave me. I'd like to acknowledge Dan Nelson for his warm friendship and support. Jerry Hasting has always been a great role model, and I appreciate his support (and his sense of humor) over the years. I'd also like to acknowledge the following colleagues from the Boy Scouts of America: Nate Langston, Rob Page, Lance Parkes,

Myra Riggs, Debra Kendrick, Carla Wallace, and Carol Barnes, all of whom have helped make my stay there one I'll always cherish. My special thanks to my friend Dabir Haider for helping with various issues time and again. I appreciate the enormous help provided by Letica Salazar over the last year.

As with all my prior books, my friend Mark Potts has helped me immensely by troubleshooting and fixing problems with my computer as well as helping with several other things. Thanks, Mark, for all the help! I'd also like to acknowledge the friendship and support of Sabrina Kirkpatrick and Al Wilson during my tenure at the Boy Scouts.

I've been working at ERCOT for only a short period, but I've benefited from the warm atmosphere and the challenging work environment. My thanks to Shawna Jirasek and Bryan Hanley for bringing me here and for their understanding and support. Thanks to the trio of exceptionally good Oracle DBAs I'm privileged to work with every day: Sudhir Mahableshwarkar, Ben Groenewald, and Bhavesh Rathod. Friends Durga Viswanatha Raju Gadiraju, Sankara Krishnaswamy, Jai Kanuri, and and Nannapaneni Sandeep also have indirectly contributed to the book by helping out in many ways in settling in at ERCOT, and I thank them for that.

I'd like to take the opportunity to thank, from the bottom of my heart, the selfless help given me by Peter Ronald, Pamela Kruger, and Angela Shoup.

My father, Appa Rao, and my mother, Swarna Kumari, have always been a source of strength and joy, and as with all my books, they have provided the inspiration for me to write. I owe my academic and professional achievements to the support and encouragement my parents have provided me over the years. My brothers, Hari and Siva Sankar, although they live far me, are close to me in spirit and inspire me every day. I'm grateful to my sisters-in-law, Aruna and Vanaja, for their support and affection. I cherish deeply the love and affection of my nephews, Teja and Ashwin, and my nieces, Aparna and Soumya. Last but not the least, I appreciate the kindness and love shown by Sobbanna Boppanna, whose well wishes have been a source of support.

Finally, I acknowledge the enormous debt I owe to all the members of my immediate family: my wife, Valerie, and children, Shannon, Nina, and Nicholas, since it's really they who have enabled me to write all of my books, by cheerfully taking care of the numerous things that fall on the wayside when one embarks on a book-writing endeavor. The children have forgiven my frequent absences (figuratively) from home as I immersed myself in this book. Without their help, love, and support, I simply couldn't have written this book. Yet again, Valerie has supported me by making numerous sacrifices so I can write. I thank her for practicing the greatest virtue of all—patience—during the long time I spent working on this book.

Introduction

GRATIANO: . . . *As who should say "I am Sir Oracle, And when I ope my lips, let no dog bark!"*

—*The Merchant of Venice*, act 1, scene 1

Oracle Corporation used to print the preceding quotation from Shakespeare at the beginning of one of its chapters in the database administrator (DBA) manual for an earlier release of the Oracle database (Oracle version 6). I always thought the quote was interesting. If you proceed a little further in the play, you'll find this quotation:

BASSANIO: *Gratiano speaks an infinite deal of nothing, more than any man in all Venice. His reasons are as two grains of wheat hid in two bushels of chaff: you shall seek all day ere you find them . . .*

—*The Merchant of Venice*, act 1, scene 1

Bassanio counters that, in truth, Gratiano speaks too much: from two bushels of chaff, two grains of wheat may be recovered. And that's the *raison d'être* for this book: to separate the wheat from the chaff. The second quotation is even more apt when you consider the difficulty of extracting the right database management procedures from the tons of material available for the Oracle Database 11g database. Oracle Corporation publishes copious material to help you manage its increasingly complex databases. Oracle Corporation also conducts a variety of in-person and web-based classes to explain the vast amount of subject matter that you need to understand to effectively work with the Oracle database today. Yet users will have a good deal of difficulty finding the essential material for performing their jobs if they rely exclusively on Oracle's voluminous (albeit well-written) material in the form of manuals, class notes, web-based seminars, and so on.

The goal of this book is to provide you with a single source for most of your day-to-day Oracle database management tasks. Of course, it isn't feasible to cover each and every DBA topic in detail. What I've done in this book is focus on the topics that are common to most enterprises, such as installing the Oracle Database 11g software, creating and upgrading databases, exporting and importing data, backing up and recovering data, and performance tuning. I place a lot of emphasis in this book on explaining all of Oracle's automatic management solutions. Using Oracle's automatic management features will keep you from reinventing the wheel each time. It also turns out that after several years of development, Oracle has finally placed in your hands a set of powerful management advisors and other tools that make a lot of traditional DBA work obsolete.

How to Become an Oracle DBA

As you start out on your journey to become a proficient Oracle DBA, you have many sources of information on the Oracle database:

- Oracle Database 11g database administration classes, which have now been boiled down to a pair of five-day long classes
- Oracle manuals—an entire library of which are available on the Oracle web sites
- Books from various publishers that impart the various pieces of knowledge you need to become an accomplished Oracle DBA

You'll also need to acquire the necessary operating system knowledge. Most of the large Oracle databases are based on the UNIX (or Linux) operating system, so you'll need to have a reasonably good understanding of UNIX. Again, you have many sources of information available. You can attend a class or two from the leading UNIX system vendors, such as Hewlett-Packard and Sun Microsystems, you can read the manuals, or you can buy some books. Microsoft Windows is another popular operating system for Oracle databases, so you need to have a basic understanding of the Windows Server operating system as well.

As many of the new entrants to the Oracle Database 11g field find out, the Oracle DBA world is exhilarating, but alas, it's also exhaustive in its reach and scope. It isn't uncommon for DBAs to have an entire shelf full of books, all explaining various facets of the DBA profession—modeling books, UNIX texts, DBA handbooks, backup and recovery guides, performance-tuning manuals, and networking and troubleshooting books. The amazing thing is, even after you run through the whole gauntlet of courses and books, you aren't really assured of being fully prepared to handle complex, day-to-day database administration chores. There are many people who have taken all the requisite classes to become an Oracle DBA who won't or can't be competent Oracle DBAs based solely on their training. The reason? Refer back to that quotation from Shakespeare at the beginning of this introduction: you need to separate the grain from the chaff, and all the coursework and manuals, while excellent in their content, can serve to muddy the waters further.

The experienced Oracle DBA can find his or her way through this baffling amount of material, but how's the neophyte DBA to cope with the overwhelming amount of information? That's where this book comes in. This text will not only educate you in the theory and principles involved in managing relational databases, it will also help you translate that theory into the useful, practical knowledge that will enable you to manage real-life Oracle Database 11g databases with real-life data and real-life issues.

Oracle Database 11g

The *g* in Oracle Database 11*g* stands for “grid.” The idea is to enable software to access spare processing power across networks (grids) of inexpensive servers. Traditionally, database systems have been run on large servers capable of running several very large databases at once. However, there are distinct disadvantages inherent in the single-server model. For example, resources tied up in the large servers can't be redistributed among the various databases and other services to ensure an optimal allocation of resources. If you need a massive amount of resources to handle your database's peak needs, chances are that you'll run with identical resources throughout the day, thus guaranteeing that you are going to waste critical resources during low-utilization periods.

Grid computing provides a means of harnessing the power of a large number of cheaper servers to supply the computing power you need in a flexible manner. This hardware would be servers like the Intel-based blade servers, and the software would include the free (or almost free) open source Linux operating system. By choosing small, generic servers, your system will cost much less than a

traditional large server system, and because you can dynamically reallocate or provision resources based on actual needs, you'll be using resources efficiently.

Grid computing (also referred to as *computing on demand* and *utility computing*) isn't a new innovation invented solely by Oracle. The idea of grid computing has been around for a while, primarily in the academic world. In fact, grid computing arose out of the academic community's need for extremely fast and scalable computers to perform complex, massive research tasks. Another overriding goal of the academic community was to permit the sharing of computing resources among large numbers of researchers. Of course, the academics also aimed to keep the cost as low as possible. Grid computing emerged out of these efforts as a viable way to create huge sharable computing environments that are dynamically adjustable to changes in the demand for computing power.

When we talk about harnessing the power of a number of commodity servers, realize that the number of computers may not be limited to just a handful. We are talking about combining the power of a fairly large number of small servers linked together to form a grid. Obviously, the key idea here is that the sum is far greater than the individual components. Enterprise grid computing, as envisioned by Oracle, uses large pools of modular storage and commodity servers. Underutilization of resources will be cut down, because capacity could be altered from the centralized pool of resources as necessary.

Here is a summary of the key benefits of grid computing:

- *Flexibility*: Since you are creating a single logical entity from a bunch of small servers, you can, of course, add or remove individual components as your computing needs dictate.
- *Efficiency*: The concept of *dynamic provisioning* underlies grid computing. Dynamic provisioning means that the allocation of resources for various services is not rigidly fixed, but changes according to the need for resources and the availability of the resources. Ideally, a well-run grid will channel resources to where they are needed the most by diverting them from underutilized sources.
- *Easy manageability*: It is far easier to manage a single logical combination of your computing resources (which may include several databases and application servers), rather than monitoring each one as a completely independent unit.
- *Economy*: The total cost of a grid environment could be considerably lower than a traditional single, big server environment. Oracle strongly recommends the use of Linux-based commodity servers, which Oracle says offer the best price/performance ratio.

Key Components of Oracle Database 11g

Following are the essential components of Oracle's grid-based systems:

- Real Application Clusters (RAC)
- Information sharing
- Easy server manageability
- Extensive instrumentation
- The advisory framework
- Automatic performance tuning
- Automatic Storage Management (ASM)
- Automatic memory management
- Scheduling and resource management
- Real Application Testing

Note that you most certainly don't have to use a "grid" platform to be able to use the Oracle Database 11g server. In either case, you can take advantage of all the new features of the database system.

Real Application Clusters

Oracle has had a feature called the Oracle Parallel Server (OPS) for many years, which enabled people to access the database from more than one instance, thus providing for scalability as well as high availability. Oracle has refined the parallel server technology considerably over the years, eventually renaming it Real Application Clusters (RAC) a few years ago.

Note This book concerns itself exclusively with the "mainstream" Oracle Database 11g DBA concepts and techniques. You'll not find any discussion of the Oracle Real Application Clusters in this book. If you are interested in RAC, you may want to take a look at Oracle manuals or refer to one of the many good books devoted to RAC.

Information Sharing

In order to efficiently share information over a grid spanning many heterogeneous systems, you need to share information efficiently. Data exchange can be occasional (such as when you perform data loads for a new system), or it could be regular and instantaneous (updating one part of the system when something changes in another part). In order to facilitate either type of information sharing, Oracle Database 11g provides *transportable tablespaces* and *Oracle Streams*.

Transportable Tablespaces

The transportable tablespaces feature enables high-speed transport of huge amounts of data from one database to another, even if the databases are running on different operating systems. The ability to move huge amounts of data across platforms, and even to rename the tablespaces during the process, makes information exchange far easier.

Oracle Streams

Oracle Streams is a feature that enables you to effortlessly capture changes made in one database and propagate them to subscriber nodes in the grid. The Oracle Streams feature can keep all the copies in sync while the changes are being applied.

Easy Server Manageability

Through its Database Control and Grid Control interfaces, Oracle Enterprise Manager enables the management of either a single database or all databases, application servers, hosts, listeners, HTTP servers, and web applications as well.

The prevailing view among IT organizations is that Oracle is a complex, difficult-to-manage database, especially when compared with the Windows server database, SQL Server. Oracle Database 11g makes a conscious effort to simplify management, right from the installation process through to daily monitoring and performance tuning. There is a new common infrastructure for storing workload- and performance-related information. You can now use powerful SQL tuning tools to determine ways to improve performance.

Oracle Enterprise Manager (OEM), which includes the single database-level Database Control, and its enterprise-wide counterpart, the Grid Control, provide unsurpassed capabilities for managing the database. Traditionally, Oracle DBAs relied on complex SQL scripts to monitor the database as well as diagnose and fix performance problems. OEM now can help you do all those things and a lot

more, without having to spend enormous amounts of time writing lengthy scripts to help manage the database.

Note I've reduced the use of DBA scripts to the bare minimum in this book. Instead, I show you how to use the OEM Database Control effectively to perform all your tasks quickly and with far less effort.

Extensive Instrumentation

Oracle Database 11g provides instrumentation of its code base that ranges further than any prior release of Oracle, providing accurate metrics about database performance that weren't available until now. Oracle's own instrumentation and metrics, since they are embedded in the database code, provide better information without any measurable performance degradation, compared to third-party performance-measurement tools.

The Advisory Framework

Oracle Database 11g contains several highly useful *advisors* to help you optimize the performance of the various components of the database. Here are some of them:

- The *Automatic Database Diagnostic Monitor* (ADDM) helps you analyze current and past instance performance.
- The *SQL Tuning Advisor* helps you tune SQL statements.
- The *SQL Access Advisor* tells you whether you should add (or drop) indexes and materialized views.
- The *Segment Advisor* helps you figure out the necessary space for new tables and to reclaim unused space assigned to segments, among other things.
- The *Undo Advisor* helps you configure the critical undo tablespace.
- The *Memory Advisor* provides recommendations for memory-related parameters.

Each of these advisors has a similar look and feel, and this consistency will help you learn how to use them effectively. Using the advisors isn't mandatory, of course—you can also tune space and memory by using Oracle-supplied packages and various dynamic performance views—but it's more efficient to simply invoke the necessary advisor.

Automatic Performance Tuning

Oracle Database 11g provides you with automatic performance diagnosis and tuning recommendations. An expert diagnosis tool called the Automatic Database Diagnostic Monitor uses the new Automatic Workload Repository contents to analyze instance performance. The ADDM's analysis includes a summary of database problems ranked according to the amount of database time they're costing, as well as a list of recommendations to eliminate these problems. The ADDM's recommendations may include modifying configuration settings or running one of the advisors listed in the previous section.

Automatic Storage Management

A significant component of the Oracle's push toward easier management is the Automatic Storage Management feature. Traditionally, database administrators relied on third-party vendors, such as VERITAS and EMC, to provide storage management tools for larger systems. ASM enables the automatic management of disks without resorting to third-party Logical Volume Managers (LVMs).

You can use Oracle's storage virtualization layer to automate and simplify the layout and management of all Oracle database files, when you use ASM. Instead of directly managing numerous files and disks, you can pay attention to a relatively small number of *disk groups*. If you need additional storage, you simply add new physical disks to the logical disk groups.

Automatic Memory Management

The Oracle Database 11g server provides you with an easy way of managing the memory needs of your databases. Automatic memory management and automatic program global area management use information collected from the instance to efficiently allocate both the major components of Oracle's memory allocation—the system global area (SGA) and the program global area (PGA).

Scheduling and Resource Management

It's common for enterprise users to share computing resources, and there needs to be a way of scheduling the users and sharing the enterprise's resources efficiently. Oracle Database 11g DBAs can use the *Database Resource Manager* to control and channel scarce database resources among the various users of the grid. You can also use the *Oracle Scheduler* to manage and monitor jobs as well as prioritize them.

Real Application Testing

Two major features of Oracle Database 11g—Database Replay and the SQL Performance Analyzer—facilitate change management by letting you replay database activities and SQL workloads, respectively. You can thus test the impact of a potential database or server upgrade, for example, by invoking the Database Replay and the SQL Performance Analyzer tools.

Why Read *This Book*?

What sets this book apart from the others on the market is the constant focus on the practical side of the DBA's work life. What does a new DBA need to know to begin work? How much and what SQL does the new DBA need to know? What UNIX, Linux, and Windows commands and utilities does the new DBA need to know? How does a DBA perform the basic UNIX administration tasks? How does a DBA install the Oracle software from scratch? How does the DBA use all the powerful new performance tuning features of the Oracle Database 11g server?

This book provides the conceptual background and operational details for all the topics a professional Oracle DBA needs to know. The following sections outline other reasons to choose this book over its competitors.

Delivers a One-Volume Reference

This book's specific purpose is to serve as a one-volume handbook for professional Oracle DBAs—as a book that covers both the theory and practice of the DBA craft. As I mentioned before, most newcomers to the field are intimidated and bewildered by the sheer amount of material they're confronted with and the great number of administrative commands they need to have at their fingertips. Well, everything you need to know to run your databases efficiently is right here in this one book.

How did I manage to achieve the difficult feat of providing comprehensive instruction in just one book? Well, although there *is* a lot of terrain to cover if you want to learn all the DBA material, you

must learn to separate the critical from the mundane, so you can identify what matters most and what you merely need to be aware of, at least in the beginning.

I'm definitely not suggesting that this one book will supplant all of the other Oracle material available. I strongly recommend that inquisitive readers make it a habit to refer to Oracle's documentation for the 11g database. You can obtain this documentation on the Web by getting a free membership to the Oracle Technology Network (OTN), which you can access through the Oracle web site at <http://technet.oracle.com>.

It's extremely important to read the Oracle database manuals and to understand how the database works. However, nothing can replace working on an actual database when it comes to mastering DBA techniques, so if you have a Windows desktop, you can easily install the freely downloadable Oracle Database 11g software. If you want, you can do the same on a Linux system as well. One of the great things about the Oracle database software is that it runs virtually identically on each operating system. In fact, your production system will operate exactly the same as the free "toy database" on your desktop machine, so go ahead and practice to your heart's content on the Oracle Database 11g database.

READING THE ORACLE MANUALS

Whether you use this or some other DBA handbook, you will still need to refer to the Oracle database manuals frequently to get the full details of complex database operations. I can't overemphasize the importance of mastering the fundamentals of Oracle Database 11g that are presented in the *Oracle Concepts* manual. Mastering this volume is critical to understanding many advanced DBA procedures.

The Oracle manuals are invaluable if you need a lot of detail. For example, the chapters on backup and recovery (Chapters 15 and 16) are good starting points in your attempt to master the Oracle procedures in those areas. Oracle has several manuals covering the backup and recovery material. Once you finish the two relevant chapters in this book, you'll find going through those manuals a pretty easy task, because you'll already have a good understanding of all the important concepts. This book provides a foundation on which you can build using the Oracle manuals and other online help available from Oracle.

In addition to the online manuals, Oracle provides an excellent set of tutorials that contain systematic instructions on how to perform many useful Oracle Database 11g tasks. You can access these tutorials, the *Oracle by Example* series, by going to <http://www.oracle.com/technology/obe/start/index.html>.

Emphasizes New Methods and When to Use Them

One of the fundamental difficulties for a neophyte in this field is determining the right strategy for managing databases. Although the essential tasks of database management are pretty similar in Oracle Database 11g compared to earlier versions of the software, the database contains several innovative techniques that make a number of routine tasks easier to perform than in the past. Oracle Corporation, however, has shied away from firmly recommending the adoption of the new methods and techniques to manage databases. The reason for this is twofold. First, Oracle rarely discards existing techniques abruptly between versions; features advertised as being destined for obsolescence are made obsolete only after many years. Thus, old and new ways of performing similar tasks coexist in the same version. Second, Oracle isn't very effective in clearly communicating its guidelines concerning contending methods. Thus, when more than one method exists for performing a task, you as a DBA have to exercise caution when you select the appropriate methods to use.

In this book, I clearly emphasize the newer features of Oracle that have been refined in the last few years and encourage you to move away from older techniques when the new innovations are clearly superior. I help you in formulating a solid strategy when multiple choices are offered. A good example is performance tuning: it was common to employ a traditional SQL-script approach to guide performance tuning efforts, but this book comes down squarely on the side of using the latest Oracle Enterprise Manager (OEM) GUI techniques to perform all your performance tuning and other DBA tasks.

Covers UNIX, SQL, PL/SQL, and Data Modeling

Some people who are motivated to become Oracle DBAs are stymied in their initial efforts to do so by their lack of training in UNIX/Linux and SQL. Also, sometimes DBAs are confused by the whole set of data modeling and the “logical DBA” techniques. This book is unique in that it covers all the essential UNIX, SQL, PL/SQL, and data modeling that a DBA ought to know to perform his or her job well.

As a DBA, you need to be able to use a number of UNIX tools and utilities to administer an Oracle database. Unfortunately, up until now many books haven’t included coverage of these vital tools. This book remedies this neglect by covering tools such as telnet, FTP, and the crontab. Many developers and managers want to have a better understanding of the UNIX system, including the use of the vi file editor, file manipulation, and basic shell-script writing techniques. This book enables you to start using the UNIX operating system right away and shows you how to write solid shell scripts to perform various tasks. Of course, you can take a specialized class or study a separate book in each of the previous areas, but that’s exactly what you’re trying to avoid by using this book.

In addition to learning all the UNIX you need to start working with the UNIX operating system right away, you can get a good working knowledge of SQL and PL/SQL from a DBA’s perspective in this book. Of course, I strongly recommend further study of both UNIX and SQL to strengthen your skills as an Oracle DBA as you progress in your career.

Offers Hands-On Administrative Experience

Although a number of books have been published in the last decade on the subject of Oracle database administration, there has been a surprising lack of the blending of the concepts of the Oracle database with the techniques needed to perform several administrative tasks. A glaring example is the area of backup and recovery, where it’s difficult to find discussions of the conceptual underpinnings of Oracle’s backup and recovery process. Consequently, many DBAs end up learning backup and recovery techniques without having a solid grasp of the underlying principles of backup and recovery. As you can imagine, this split between theory and practice proves expensive in the middle of a recovery operation, where fuzziness on the concepts could lead to simple mistakes.

Your success as a professional database administrator is directly related to the amount of hands-on experience you have, and to your understanding of the concepts behind the operation of the database. To get this practice, you can experiment with all the commands in this book on a UNIX- or a Windows-based Oracle Database 11g database. Oracle Database 11g is loaded with features that make it the cutting-edge database in the relational database market, and this book covers all the new additions and modifications to database administration contained in the 11g version. It’s a lot of fun for an experienced DBA to have the opportunity to use all the wonderful features of the new database, but beginning- and intermediate-level DBAs will have *more* fun, because they’re embarking on the great endeavor that is the mastery of Oracle database management.

Who Should Read This Book?

This book is primarily intended for beginning- and intermediate-level Oracle Database 11g DBAs. Prior experience with Oracle databases isn't assumed, so if you've never managed databases and intend to master the management of the new Oracle Database 11g database, you can do so with the help of this book.

More precisely, the audience for this book will fall into the following categories:

- Oracle DBAs who are just starting out
- Oracle developers and UNIX/Linux or Windows system administrators who intend to learn Oracle DBA skills
- Managers who want to get a hands-on feel for database management
- Anybody who wants to learn how to become a proficient Oracle DBA on his or her own

A Note About UNIX, Linux, and Windows

I personally like the UNIX operating system and use it at work. I'm familiar with the Windows platform and I think it's a good operating system for small enterprises, but my favorite operating system remains UNIX, which stands out for its reliability, scalability, and speed. For medium and large organizations, the UNIX system offers wonderful features and ease of use. As a result, you'll find this book heavily oriented toward the use of Oracle on UNIX systems.

If you happen to admire the Linux operating system, there isn't a new learning curve involved, as most of the operating system commands will work the same way in the UNIX and Linux systems. If you need to find out how to use the Oracle Database 11g database on a Windows platform, here's some interesting news for you: the database commands work exactly the same way in both the UNIX and Windows environments.

How This Book Is Organized

I have organized the contents of this book with the new DBA in mind. My goal is to provide you with a decent background in data modeling, SQL, and UNIX, while providing a thorough course in the essentials of Oracle Database 11g database management skills. I know it's unusual to provide UNIX and SQL background in an Oracle DBA book, but this inclusion is in line with the goal I set when I decided to write this book: there ought to be a single book or manual that has all the necessary background for a reader to start working as an Oracle Database 11g DBA.

I strove to write the chapters to mirror real-life practical training. For example, you should understand basic database modeling and fundamental UNIX operating system commands before learning to manage Oracle databases. I therefore start with a discussion of database modeling and UNIX. You'll install the Oracle database software before learning how to create an Oracle database. After you install the software and create a database, you can create users and establish connectivity. Subsequent chapters deal with the loading and unloading of data, backup and recovery, day-to-day database management, and performance tuning.

I advise beginning DBAs to start at the beginning of the book and keep going. A more experienced user, on the other hand, can pick the topics in any sequence he or she desires. Throughout the book, I've provided detailed, step-by-step, tested examples to illustrate the use of data concepts and

features of Oracle Database 11g. I strongly recommend that you set up an Oracle Database 11g database server on your PC and follow along with these examples. Doing so will teach you the relevant commands and help you build confidence in your skill level. Moreover, the examples are a whole lot of fun! The following sections briefly summarize the contents of the book.

Part 1: Background, Data Modeling, UNIX/Linux, and SQL*Plus

Part 1 provides a background on the Oracle DBA profession and offers an introduction to data modeling and the UNIX operating system as well as SQL*Plus. In Chapter 1 I discuss the role of the Oracle DBA in the organization, and I offer some advice on improving your skill set as a DBA. I also discuss the basics of relational databases. Chapter 2 provides an introduction to both logical and physical database design, including the use of entity-relationship diagrams. You'll learn about the Optimal Flexible Architecture (OFA) with regard to disk layout. Chapter 3 provides a quick introduction to UNIX/Linux operating systems, including the most common commands that you need as an Oracle DBA, the rudiments of shell scripting, and how to use the vi text-processing commands. You'll also explore the essential UNIX system administration tasks for Oracle DBAs. This chapter finishes with coverage of disks and storage systems, including the popular RAID systems. Chapter 4 provides a thorough introduction to the use of SQL* Plus, the main interface to the Oracle database. In addition, Chapter 4 also describes how to use the powerful Oracle Enterprise Manager to monitor and manage your databases as well as your entire system. You'll learn how to install and use the Database Control, which you use for managing a single database, and the Grid Control, through which you can manage your enterprise, including application servers and hosts.

Part 2: Oracle Database 11g Architecture, Schema, and Transaction Management

Part 2 is in many ways the heart of the book—it covers the important topics of Oracle Database 11g's architecture, schema management, and transaction management. In Chapter 5 you'll learn about the important components of the Oracle database architecture, such as how the database processes and memory work. It also covers the conceptual foundations of the Oracle database. Chapter 6 provides a detailed introduction to the management of tablespaces. Chapter 7 covers schema management in Oracle Database 11g, and it contains a quick review of the important types of Oracle objects, such as tables and indexes, and shows you how to manage them. Chapter 8 provides you with a thorough understanding of how Oracle databases conduct transaction processing.

Part 3: Installing Oracle Database 11g, Upgrading, and Creating Databases

Part 3 includes two chapters that show you how to install the Oracle Database 11g software, create Oracle databases, and upgrade databases. Chapter 9, which covers Oracle software installation, shows how to install the Oracle Database 11g database server. In addition, Chapter 9 also shows you in detail how to upgrade to Oracle Database 11g. Chapter 10 shows you how to create an Oracle database from scratch, both manually as well as by using the Database Configuration Assistant (DBCA).

Part 4: Connectivity and User Management

Part 4 explains how to establish connectivity to the Oracle database and manage database users. Chapter 11 covers connecting to Oracle databases, and Chapter 12 shows you how to manage users and discusses ways of securing your database.

Part 5: Data Loading, Backup, and Recovery

Part 5 deals with loading data and performing backups and recovery. You'll learn how to use SQL*Loader in Chapter 13, and Chapter 14 covers the Data Pump technology, which enables you to load and unload Oracle data. Chapters 15 and 16 deal with the crucial topics of database backups and recovery, respectively.

Part 6: Managing the Database

Part 6 covers managing the operational Oracle Database 11g database. Chapter 17 focuses on the important Oracle Database 11g automatic management features, as well as exploring several powerful online capabilities of the Oracle database. Chapter 18 shows you how to manage data files, tablespaces, and Oracle redo logs, and how to perform undo management. The chapter also provides an introduction to the Oracle storage solution, Automatic Storage Management.

Part 7: Performance Tuning

Part 7 covers Oracle Database 11g performance tuning and troubleshooting issues. Chapter 19 discusses the Cost-Based Optimizer and provides tips on writing efficient SQL queries. You'll also see how to use Oracle's Automatic SQL Tuning Advisor to improve query performance. In Chapter 20, you'll learn how to optimize the use of Oracle's memory, disk I/O, and the operating system. You'll also learn about the Oracle wait interface in this chapter. A basic approach to performance analysis and troubleshooting production databases is explained as well.

Appendix: Oracle Database 11g SQL and PL/SQL: A Brief Primer

In the Appendix, I introduce Oracle SQL and PL/SQL, provide an introduction to Oracle XML DB, which helps you deal with XML data, and include an introduction to using the Java programming language with Oracle.

Salud!

I truly enjoy working with the Oracle database, because of its amazing range of capabilities and the intricate challenges it throws my way as I explore its wide-ranging capabilities. I hope you derive as much satisfaction and fulfillment from the Oracle database as I do. I leave you with the following observation, adapted from the introduction to the famous economics textbook by Paul A. Samuelson, the great economist and Nobel Laureate:¹

I envy you, the beginning Oracle DBA, as you set out to explore the exciting world of Oracle Database 11g database management for the first time. This is a thrill that, alas, you can experience only once in a lifetime. So, as you embark, I wish you bon voyage!

1. Paul A. Samuelson and William D. Nordhaus, *Economics, Seventeenth Edition* (New York: McGraw-Hill, 1998).

